

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
from google.colab import files
upload=files.upload()
df = pd.read_csv('abalone.csv')
```

Choose Files abalone.csv

- **abalone.csv**(text/csv) - 191962 bytes, last modified: 11/5/2022 - 100% done
Saving abalone.csv to abalone.csv

```
df.describe()
```



	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	
count	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	41
mean	0.523992	0.407881	0.139516	0.828742	0.359367	0.180594	
std	0.120093	0.099240	0.041827	0.490389	0.221963	0.109614	
min	0.075000	0.055000	0.000000	0.002000	0.001000	0.000500	
25%	0.450000	0.350000	0.115000	0.441500	0.186000	0.093500	
50%	0.545000	0.425000	0.140000	0.799500	0.336000	0.171000	
75%	0.615000	0.480000	0.165000	1.153000	0.502000	0.253000	



Automatic saving failed. This file was updated remotely or in another tab.
[diff](#)

[Show](#)

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7


```
df.tail()
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
4172	F	0.565	0.450	0.165	0.8870	0.3700	0.2390	0.2490	11
4173	M	0.590	0.440	0.135	0.9660	0.4390	0.2145	0.2605	10
4174	M	0.600	0.475	0.205	1.1760	0.5255	0.2875	0.3080	9
4175	F	0.625	0.485	0.150	1.0945	0.5310	0.2610	0.2960	10

Univariate analysis

```
sns.boxplot(df.Length)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass
FutureWarning
<matplotlib.axes._subplots.AxesSubplot at 0x7fde01392090>
```




Automatic saving failed. This file was updated remotely or in another tab.
[diff](#)

[Show](#)

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fde012ebhd0>
sns.heatmap(df.isnull())
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fde00dba850>
```




▼ Bivariate analysis

```
sns.barplot(x=df.Height,y=df.Diameter)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fddfe4e7c10>
```



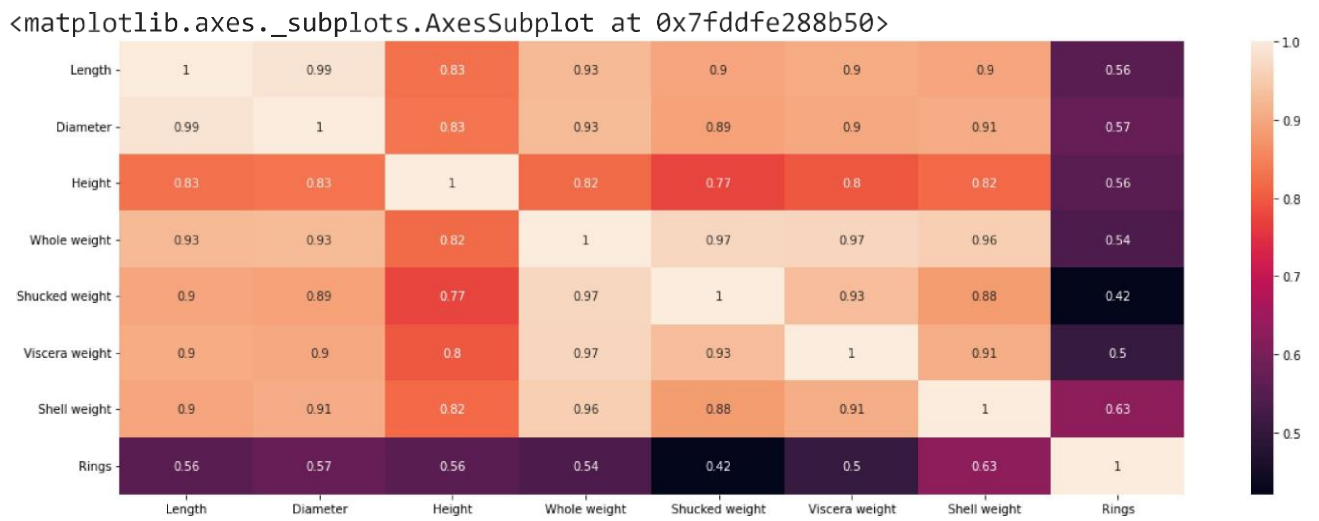
Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)



```
numerical_features = df.select_dtypes(include=[np.number]).columns
categorical_features = df.select_dtypes(include=[np.object]).columns
```

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:2: DeprecationWarning: `np.object` is deprecated in NumPy 1.20; for more details and guidance: https://numpy.org/devdocs/numpy_1_20_0_migration_guide.html

```
plt.figure(figsize = (20,7))
sns.heatmap(df[numerical_features].corr(),annot = True)
```

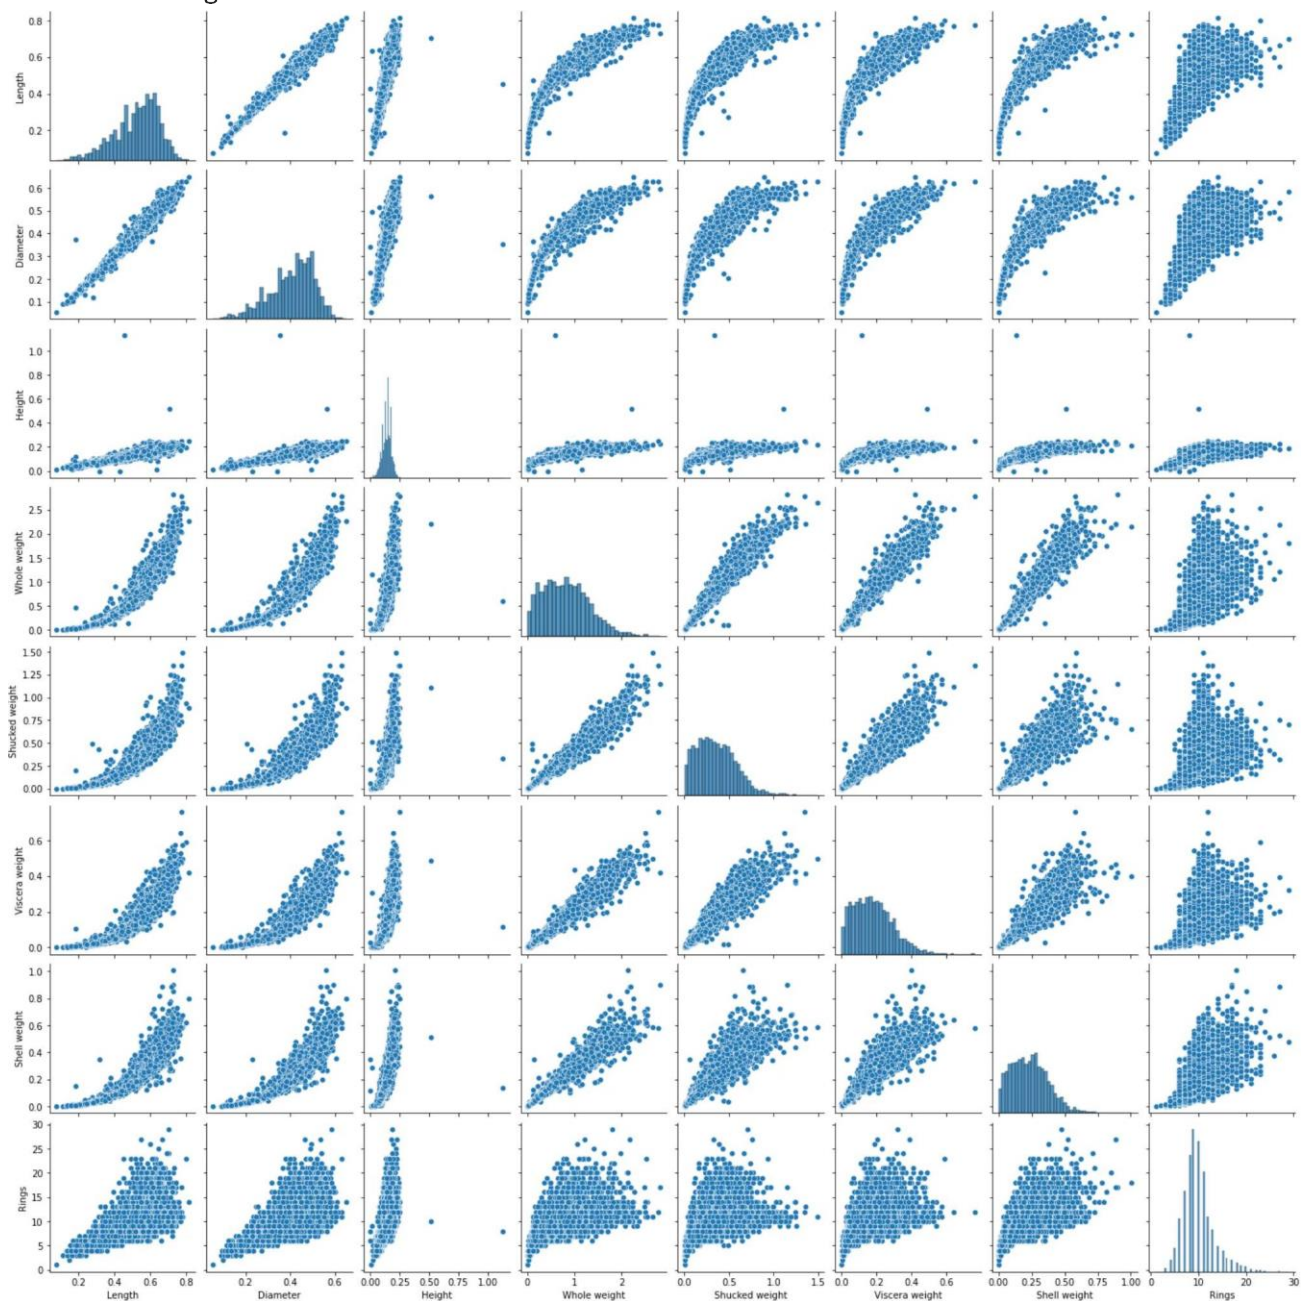


▼ Multivariate Analysis

```
sns.pairplot(df)
```

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

<seaborn.axisgrid.PairGrid at 0x7fddfe1e85d0>



Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

▼ Perform descriptive model on the dataset

```
df['Height'].describe()
```

```
count    4177.000000
mean      0.139516
std       0.041827
min       0.000000
25%      0.115000
50%      0.140000
75%      0.165000
max       1.130000
Name: Height, dtype: float64
```

```
df['Height'].mean()
```

```
0.13951639932966242
```

```
df.max()
```

```
Sex      M
Length    0.815
Diameter  0.13
Height    2.8255
Whole weight  1.488
Shucked weight  0.76
Viscera weight
Shell weight  1.005
```

Automatic saving failed. This file was updated remotely or in another tab.

[Show](#)

```
df['Sex'].value_counts()
```

```
M    1528
I    1342
F    1307
Name: Sex, dtype: int64
```

```
df[df.Height == 0]
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
--	-----	--------	----------	--------	--------------	----------------	----------------	--------------	-------

```
df['Shucked weight'].kurtosis()

0.5951236783694207

df['Diameter'].median()

0.425

df['Shucked weight'].skew()

0.7190979217612694
```

Missing values

```
df.isna().any()

Sex                False
Length             False
Diameter           False
Height             False
Whole weight       False
Shucked weight     False
Viscera weight     False
Shell weight       False
Rings              False
dtype: bool
```

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

```
percentage_missing_values = (missing_values/len(df))*100
pd.concat([missing_values, percentage_missing_values], axis = 1, keys=
```

	Missing values	% Missing	
Sex	0	0.0	
Length	0	0.0	


▼ Find the outliers

```
q1=df.Rings.quantile(0.25)
q2=df.Rings.quantile(0.75)
iqr=q2-q1
```

```
print(iqr)
```

```
3.0
```

```
df = pd.get_dummies(df)
dummy_df = df
df.boxplot( rot = 90, figsize=(20,5))
```




Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

```
df['age'] = df['Rings']
df = df.drop('Rings', axis = 1)
```

```
df.drop(df[(df['Viscera weight'] > 0.5) & (df['age'] < 20)].index, inplace=True)
df.drop(df[(df['Viscera weight'] < 0.5) & (df['age'] > 25)].index, inplace=True)
```



```
var = 'Shell weight'
plt.scatter(x = df[var], y = df['age'])
plt.grid(True)
```



▼ Check for categorical columns and perform encoding

```
numerical_features = df.select_dtypes(include = [np.number]).columns
categorical_features = df.select_dtypes(include = [np.object]).columns
```

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:2: DeprecationWarning: `DeprecationWarning: ` is deprecated in NumPy 1.20; for more details and guidance: <https://numpy.org/devdocs/numpy-2.0-0-notes>

```
numerical_features
categorical_features
```

```
Index([], dtype='object')
```

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

```
abalone_numeric.head()
```

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	age	Sex_F	Sex_I	Sex_M
0	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15	0	0	0
1	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7	0	0	0
2	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9	1	0	0
3	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10	0	0	0

▼ Dependent and Independent Variables

```
x = df.iloc[:, 0:1].values
y = df.iloc[:, 1]
y
```

```
0      0.365
1      0.265
2      0.420
3      0.365
4      0.255
...
4172   0.450
4173   0.440
4174   0.475
4175   0.485
4176   0.555
Name: Diameter, Length: 4150, dtype: float64
```

```
#Scaling the Independent Variables
print ("\n ORIGINAL VALUES: \n\n", x,y)
```

ORIGINAL VALUES:

```
[[0.455]
 [0.35 ]
 [0.53 ]
 ...
 [0.6  ]
 [0.625]
 [0.71 ]] 0      0.365
1      0.265
2      0.420
3      0.365
```

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

```
4174   0.475
4175   0.440
4176   0.555
Name: Diameter, Length: 4150, dtype: float64
```

```
import preprocessing
min_max_scaler = preprocessing.MinMaxScaler(feature_range =(0, 1))
new_y = min_max_scaler.fit_transform(x,y)
print ("\n VALUES AFTER MIN MAX SCALING: \n\n", new_y)
```

VALUES AFTER MIN MAX SCALING:

```
[[0.51351351]
 [0.37162162]
 [0.61486486]
 ...
 [0.70945946]
 [0.74324324]
 [0.85810811]]
```

#Split the data into Training and Testing

```
X = df.drop('age', axis = 1)
```

```
Y = df['age']
```

```
from sklearn.preprocessing import StandardScaler
```

```
from sklearn.model_selection import train_test_split, cross_val_score
```

```
from sklearn.feature_selection import SelectKBest
```

```
standardScale = StandardScaler()
```

```
standardScale.fit_transform(X)
```

```
selectkBest = SelectKBest()
```

```
X_new = selectkBest.fit_transform(X, y)
```

```
X_train, X_test, y_train, y_test = train_test_split(X_new, y, test_size = 0.25)
```

```
X_train
```

```
array([[0.505, 0.39 , 0.12 , ..., 1.    , 0.    , 0.    ],
       [0.69 ,
       [0.27 , 0.55 , 0.18 , ..., 0.    , 0.    , 1.    ],
       ...,    0.195, 0.07 , ..., 0.    , 0.    , 1.    ],
       [0.67 , 0.51 ,
       [0.325,    0.155, ..., 1.    , 0.    , 0.    ],
       [0.41 , 0.24 , 0.075, ..., 0.    , 1.    , 0.    ],
       0.325, 0.1  , ..., 0.    , 1.    , 0.    ]])
```

```
y_train
```

```
3447      8
1975     11
```

```
2149      7
```

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

```
1407     18
```

```
3703     11
```

```
3430      6
```

```
1075      6
```

```
Name: age, Length: 3112, dtype: int64
```

Build the model

Linear Regression

```
from sklearn import linear_model as lm
```

```
from sklearn.linear_model import LinearRegression
model=lm.LinearRegression()
results=model.fit(X_train,y_train)
```

```
accuracy = model.score(X_train, y_train)
print('Accuracy of the model:', accuracy)
```

Accuracy of the model: 0.528142126401383

```
#Training the model
lm = LinearRegression()
lm.fit(X_train, y_train)
y_train_pred = lm.predict(X_train)
y_train_pred
```

array([9.28125, 13.90625, 7.125 , ..., 11.1875 , 6.65625, 8.0625])

X_train

```
array([[0.505, 0.39 , 0.12 , ..., 1. , 0. , 0. ],
       [0.69 , 0.55 , 0.18 , ..., 0. , 0. , 1. ],
       [0.27 , 0.195, 0.07 , ..., 0. , 0. , 1. ],
       ...,
       [0.67 , 0.51 , 0.155, ..., 1. , 0. , 0. ],
       [0.325, 0.24 , 0.075, ..., 0. , 1. , 0. ],
       [0.41 , 0.325, 0.1 , ..., 0. , 1. , 0. ]])
```

y_train

```
3447      8
1975     11
2149      7
2678      7
3301     13
3403     18
3430      6
1075      6
```

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

```
3403     18
3430      6
1075      6
Name: age, Length: 3112, dtype: int64
```

```
sklearn.metrics import mean_absolute_error, mean_squared_error
frommean_squared_error(y_train, y_train_pred)
print('Mean Squared error of training set :%2f'%s)
```

Mean Squared error of training set :4.933080

Testing the model

```
y_train_pred = lm.predict(X_train)
y_test_pred = lm.predict(X_test)
```

y_test_pred

```
array([16.25 , 11. , 9.25 , ..., 12.1875 , 10.53125, 5.1875 ])
```

X_test

```
array([[0.595, 0.495, 0.185, ..., 1. , 0. , 0. ],
       [0.605, 0.485, 0.16 , ..., 1. , 0. , 0. ],
       [0.52 , 0.39 , 0.12 , ..., 0. , 0. , 1. ],
       ...,
       [0.635, 0.515, 0.165, ..., 0. , 0. , 1. ],
       [0.565, 0.45 , 0.175, ..., 1. , 0. , 0. ],
       [0.2 , 0.145, 0.025, ..., 0. , 1. , 0. ]])
```

y_test

```
67      13
161     13
3448      7
4019     10
378     15
      ..
984     10
3862    10
1948    10
1132      9
3190      5
Name: age, Length: 1038, dtype: int64
```

Automatic saving failed. This file was updated remotely or in another tab. [Show](#)

[diff](#) Mean Squared error of testing set :4.058311

Measure the performance using metrics

```
sklearn.metrics import r2_score
from sklearn.metrics import r2_score
print('R2 Score of training set: %.2f' % r2_score(y_train, y_train_pred))
```

```
R2 Score of training set:0.53
```

```
from sklearn.metrics import r2_score
```

```
p = r2_score(y_test, y_test_pred)
print('R2 Score of testing set: %.2f'%p)
```

```
R2 Score of testing set: 0.56
```

[Colab paid products](#) - [Cancel contracts here](#)

✓ 0s completed at 21:22